

# life\_prediction02

June 12, 2021

## 0.1 Devices Life Prediction Version 0.2

Developed by A.Okada, T.Shirakami, K.Kuramitsu, K.Iino and N.Yamazaki

Copyright© A.Okada, T.Shirakami, K.Kuramitsu, K.Iino and N.Yamazaki, 2021

Start from June 8, 2021

### 0.1.1 Sample Generator

```
[1]: # Library import
import pandas as pd # =DataFrame
import numpy as np #
import random # random
import math # log sin, cos
import matplotlib.pyplot as plt #
import matplotlib as mpl #

# DataFrame
pd.options.display.precision = 2

Sampling = 1000
NumSample = 100

# FAN
def sample_generator(RpmSpec,
                    PowerSpec,
                    TempSpec,
                    DiameterSpec,
                    ThicknessSpec,
                    QSpec,
                    PSpec,
                    LifeSpec,
                    Sampling,
                    NumSample):

    data_list = [[]] # sampling list 2 list
    # cum_rpm = 0
```

```

#     cum_temp = 0
#     cum_power = 0

for sample_id in range(NumSample):

    rpm = RpmSpec #      0
    cum_rpm = 0
    cum_temp = 0
    cum_power = 0
    cumurated_life_impact_factor = 0

    if np.random.random() < 0.1: #      10%           ==>rpm   power
        defect = 1
    else:
        defect = 0

    for time in range(0, LifeSpec*2, Sampling):

        temp = 25 + random.uniform(-5,5)

        if rpm <= 0: #   for
            power = 0
            death = 1
        else:

            if defect == 1: #
                # rpm; 40degC ,           +/-5%
                rpm = (-1 * ((time + Sampling)/8000) ** 6 + RpmSpec * temp /
→ TempSpec) * (1-random.uniform(-0.05,0.05))
                if rpm < 0.1*RpmSpec:
                    rpm = 0
                    power = 0
                    death = 1
                    remaining_life = 0
                else:
                    #power: rpm           +/-5%
                    power = (0.5 * (4000/rpm) ** 1.2 + PowerSpec * (temp /
→ TempSpec)) * (1 - random.uniform(-0.05,0.05))
                    death = 0
                    # remaining_life
                    remaining_life = ((RpmSpec*temp/TempSpec - 0.
→ 1*RpmSpec)**(1/6)) * 8000 * (1 - random.uniform(-0.05,0.05)) - time
                    if remaining_life < 0:
                        remaining_life = 0
                else: #
                    rpm = (-1 * ((time + Sampling)/8000) ** 4 + RpmSpec * temp /
→ TempSpec) * (1-random.uniform(-0.05,0.05))

```

```

        if rpm < 100:
            rpm = 0
            power = 0
            death = 1
            remaining_life = 0
        else:
            power = (0.5 * (4000/rpm) ** 1.1 + PowerSpec * (temp /
→TempSpec)) * (1 - random.uniform(-0.05,0.05))
            death = 0
            remaining_life = ((RpmSpec*temp/TempSpec - 0.
→1*RpmSpec)**(1/4)) * 8000 * (1 - random.uniform(-0.05,0.05)) -time
            if remaining_life < 0:
                remaining_life = 0

        # cum = cumurated =
        cum_rpm += rpm
        cum_temp += temp
        cum_power += power

        # FAN          rpm , temp, power
        # cumurated_life_impact_factor 0 +/-1

        cumurated_life_impact_factor = math.log(10, ((1/cum_rpm) ** 0.5) *
→cum_temp * cum_power)

        data_list.append([sample_id,
                           defect,
                           time,
                           rpm,
                           temp,
                           power,
                           cum_rpm,
                           cum_temp,
                           cum_power,
                           RpmSpec,
                           PowerSpec,
                           DiameterSpec,
                           ThicknessSpec,
                           QSpec,
                           PSpec,
                           LifeSpec,
                           cumurated_life_impact_factor,
                           death,
                           remaining_life])

    return data_list

```

```
fan40 = sample_generator(RpmSpec = 25000,
```

```

        PowerSpec = 20.16,
        TempSpec = 40,
        DiameterSpec = 40,
        ThicknessSpec = 28,
        QSpec = 0.83, #  $m^3/min$ 
        PSpec = 1100, #  $Pa$ 
        LifeSpec = 40000,
        Sampling = Sampling,
        NumSample = NumSample)

fan40cr = sample_generator(RpmSpec = 22000,
        PowerSpec = 19.2,
        TempSpec = 40,
        DiameterSpec = 40,
        ThicknessSpec = 56,
        QSpec = 0.9, #  $m^3/min$ 
        PSpec = 1045, #  $Pa$ 
        LifeSpec = 40000,
        Sampling = Sampling,
        NumSample = NumSample)

fan120 = sample_generator(RpmSpec = 7650,
        PowerSpec = 1.3 * 48,
        TempSpec = 40,
        DiameterSpec = 120,
        ThicknessSpec = 38,
        QSpec = 7.49, #  $m^3/min$ 
        PSpec = 532.5, #  $Pa$ ,
        LifeSpec = 40000,
        Sampling = Sampling,
        NumSample = NumSample)

# fan
list = fan40
for data in fan40cr:
    list.append(data)

for data in fan120:
    list.append(data)

df = pd.DataFrame(list, # list 3 fan40, fan40cr or fan120
        columns=['sample_id',
        'defect',
        'time',
        'rpm',
        'temp',
        'power',
        'cum_rpm',

```

```

        'cum_temp',
        'cum_power',
        'RpmSpec',
        'PowerSpec',
        'DiameterSpec',
        'ThicknessSpec',
        'QSpec',
        'PSpec',
        'LifeSpec',
        'cumurated_life_impact_factor',
        'death',
        'remaining_life'])

# df.to_csv("./sample_data_check3.csv")

print('df= ', df.info())

fig, ax = plt.subplots()
ax.scatter(df['time'], df['rpm'], c=df['sample_id'], s=10, alpha=0.5)
plt.xlabel('time [H]',size=12)
plt.ylabel('rpm',size=12)

fig, ax = plt.subplots()
ax.scatter(df['time'], df['power'], c=df['sample_id'])
plt.xlabel('time [H]',size=12)
plt.ylabel('power',size=12)

fig, ax = plt.subplots()
ax.scatter(df['time'], df['remaining_life'], c=df['sample_id'])
plt.xlabel('time [H]',size=12)
plt.ylabel('remaining_life [H]',size=12)

fig, ax = plt.subplots()
ax.scatter(df['time'], df['cum_rpm'], c=df['sample_id'])
plt.xlabel('time [H]',size=12)
plt.ylabel('cum_rpm',size=12)

fig, ax = plt.subplots()
ax.scatter(df['time'], df['cum_power'], c=df['sample_id'])
plt.xlabel('time [H]',size=12)
plt.ylabel('cum_power',size=12)



fig, ax = plt.subplots()
ax.scatter(df['time'], df['cum_temp'], c=df['sample_id'])
plt.xlabel('time [H]',size=12)
plt.ylabel('cum_temp',size=12)



```

```
fig, ax = plt.subplots()
ax.scatter(df['time'], df['cumurated_life_impact_factor'], c=df['sample_id'])
plt.xlabel('time [H]',size=12)
plt.ylabel('cumurated_life_impact_factor',size=12)
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 24003 entries, 0 to 24002
Data columns (total 19 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   sample_id                            24000 non-null  float64
1   defect                               24000 non-null  float64
2   time                                 24000 non-null  float64
3   rpm                                  24000 non-null  float64
4   temp                                 24000 non-null  float64
5   power                                24000 non-null  float64
6   cum_rpm                              24000 non-null  float64
7   cum_temp                             24000 non-null  float64
8   cum_power                            24000 non-null  float64
9   RpmSpec                              24000 non-null  float64
10  PowerSpec                            24000 non-null  float64
11  DiameterSpec                         24000 non-null  float64
12  ThicknessSpec                        24000 non-null  float64
13  QSpec                                24000 non-null  float64
14  PSpec                                24000 non-null  float64
15  LifeSpec                             24000 non-null  float64
16  cumurated_life_impact_factor          24000 non-null  float64
17  death                                24000 non-null  float64
18  remaining_life                        24000 non-null  float64
dtypes: float64(19)
memory usage: 3.5 MB
df= None
```


```
[1]: Text(0, 0.5, 'cumurated_life_impact_factor')
```











```
[2]: df = df.dropna(how="any")

indexNames = df[ df['death'] == 1 ].index
df.drop(indexNames , inplace=True)

indexNames = df[ df['remaining_life'] == 0 ].index
df.drop(indexNames , inplace=True)

df.to_csv('./sample_data.csv')
df
```

```
[2]:
```

	sample_id	defect	time	rpm	temp	power	cum_rpm	\
1	0.0	0.0	0.0	14399.74	22.95	11.21	14399.74	
2	0.0	0.0	1000.0	14942.86	23.66	11.88	29342.61	
3	0.0	0.0	2000.0	14045.38	22.89	11.78	43387.99	
4	0.0	0.0	3000.0	12900.79	20.52	10.32	56288.78	
5	0.0	0.0	4000.0	14429.85	23.39	12.27	70718.63	
...	...	...	...	...	...	...	...	
23981	99.0	0.0	58000.0	1247.23	22.25	36.17	247698.27	
23982	99.0	0.0	59000.0	1344.03	23.52	36.72	249042.30	
23983	99.0	0.0	60000.0	1961.37	28.01	44.69	251003.67	
23984	99.0	0.0	61000.0	644.90	22.33	36.86	251648.56	
23985	99.0	0.0	62000.0	1452.19	27.35	42.94	253100.76	

	cum_temp	cum_power	RpmSpec	PowerSpec	DiameterSpec	ThicknessSpec	\
1	22.95	11.21	25000.0	20.16	40.0	28.0	
2	46.61	23.09	25000.0	20.16	40.0	28.0	
3	69.50	34.88	25000.0	20.16	40.0	28.0	
4	90.01	45.20	25000.0	20.16	40.0	28.0	
5	113.40	57.47	25000.0	20.16	40.0	28.0	
...	...	...	...	...	...	...	
23981	1493.18	2370.01	7650.0	62.40	120.0	38.0	
23982	1516.69	2406.73	7650.0	62.40	120.0	38.0	
23983	1544.70	2451.42	7650.0	62.40	120.0	38.0	
23984	1567.03	2488.27	7650.0	62.40	120.0	38.0	
23985	1594.39	2531.21	7650.0	62.40	120.0	38.0	

	QSpec	PSpec	LifeSpec	cumurated_life_impact_factor	death	\
1	0.83	1100.0	40000.0		3.02	0.0
2	0.83	1100.0	40000.0		1.25	0.0
3	0.83	1100.0	40000.0		0.94	0.0
4	0.83	1100.0	40000.0		0.81	0.0
5	0.83	1100.0	40000.0		0.72	0.0
...	...	...	...	...	...	
23981	7.49	532.5	40000.0		0.26	0.0
23982	7.49	532.5	40000.0		0.26	0.0
23983	7.49	532.5	40000.0		0.26	0.0
23984	7.49	532.5	40000.0		0.26	0.0
23985	7.49	532.5	40000.0		0.26	0.0

	remaining_life
1	86425.74
2	85610.30
3	83857.05
4	80516.46
5	82687.82
...	...
23981	1504.73
23982	6273.60
23983	3065.98
23984	2969.42
23985	6288.92

[21041 rows x 19 columns]

### 0.1.2 Prediction

```
[3]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```


```
plt.style.use('ggplot')
import seaborn as sns

df = pd.read_csv('sample_data.csv', index_col=[0])
df = df.dropna(how="any")

# print(df.head(), df.tail())
df.info()
# print(df.describe())
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 21041 entries, 1 to 23985
Data columns (total 19 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   sample_id                            21041 non-null  float64
1   defect                               21041 non-null  float64
2   time                                 21041 non-null  float64
3   rpm                                  21041 non-null  float64
4   temp                                 21041 non-null  float64
5   power                                21041 non-null  float64
6   cum_rpm                              21041 non-null  float64
7   cum_temp                             21041 non-null  float64
8   cum_power                            21041 non-null  float64
9   RpmSpec                              21041 non-null  float64
10  PowerSpec                            21041 non-null  float64
11  DiameterSpec                         21041 non-null  float64
12  ThicknessSpec                        21041 non-null  float64
13  QSpec                                21041 non-null  float64
14  PSpec                                21041 non-null  float64
15  LifeSpec                             21041 non-null  float64
16  cumurated_life_impact_factor          21041 non-null  float64
17  death                                21041 non-null  float64
18  remaining_life                        21041 non-null  float64
dtypes: float64(19)
memory usage: 3.2 MB
```

```
[4]: plt.figure(figsize=(8, 8))
plt.hist(df['remaining_life'], bins=100)
plt.title('Data Histogram for prediction value')
plt.xlabel('remaining_life',size=12)
plt.show()
```



```
[5]: from sklearn.model_selection import train_test_split
import xgboost as xgb

#         X,         y
X = df.drop(columns=['remaining_life', 'sample_id', 'defect'])
y = df['remaining_life']

print(X.shape)
print(y.shape)

# train      test
#           20%
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=0)

print(X_train.shape)
print(X_test.shape)
```

```
(21041, 16)
(21041,)
(16832, 16)
(4209, 16)
```

```
[6]: params = {
    'silent': 1,
    'max_depth': 6,
    'min_child_weight': 1,
    'eta': 0.1,
    'tree_method': 'exact',
    'objective': 'reg:linear',
    'eval_metric': 'rmse',
    'predictor': 'cpu_predictor'
}

# GPU
# params = {
#     'silent': 1,
#     'max_depth': 6,
#     'min_child_weight': 1,
#     'eta': 0.1,
#     'tree_method': 'gpu_exact',
#     'objective': 'gpu:reg:linear',
#     'eval_metric': 'rmse',
#     'predictor': 'gpu_predictor'
# }

dtrain = xgb.DMatrix(X_train, label=y_train)
dtest = xgb.DMatrix(X_test, label=y_test)
model = xgb.train(params=params,
                  dtrain=dtrain,
                  num_boost_round=1000,
                  early_stopping_rounds=5,
                  evals=[(dtest, 'test')])
```

```
[09:07:11] WARNING: C:/Users/Administrator/workspace/xgboost-
win64_release_1.4.0/src/objective/regression_obj.cu:171: reg:linear is now
deprecated in favor of reg:squarederror.
[09:07:11] WARNING: C:/Users/Administrator/workspace/xgboost-
win64_release_1.4.0/src/learner.cc:573:
Parameters: { "silent" } might not be used.
```

This may not be accurate due to some parameters are only used in language bindings but passed down to XGBoost core. Or some parameters are not used but slip through this verification. Please open an issue if you find above cases.

```
[0]    test-rmse:41641.05859
[1]    test-rmse:37699.17969
[2]    test-rmse:34178.75391
[3]    test-rmse:31037.64648
[4]    test-rmse:28216.15430
[5]    test-rmse:25723.76758
[6]    test-rmse:23487.06055
[7]    test-rmse:21513.18359
[8]    test-rmse:19757.56250
[9]    test-rmse:18200.48047
[10]   test-rmse:16831.53711
[11]   test-rmse:15633.85352
[12]   test-rmse:14587.16113
[13]   test-rmse:13696.62305
[14]   test-rmse:12913.14648
[15]   test-rmse:12243.71387
[16]   test-rmse:11663.33301
[17]   test-rmse:11173.53906
[18]   test-rmse:10760.42383
[19]   test-rmse:10407.97461
[20]   test-rmse:10125.70508
[21]   test-rmse:9876.97754
[22]   test-rmse:9673.81543
[23]   test-rmse:9503.46191
[24]   test-rmse:9354.51172
[25]   test-rmse:9235.05664
[26]   test-rmse:9142.36914
[27]   test-rmse:9063.54785
[28]   test-rmse:8998.35938
[29]   test-rmse:8949.75684
[30]   test-rmse:8905.86914
[31]   test-rmse:8875.48047
[32]   test-rmse:8846.91406
[33]   test-rmse:8817.47266
[34]   test-rmse:8808.79102
[35]   test-rmse:8796.17481
[36]   test-rmse:8787.79199
[37]   test-rmse:8781.47266
[38]   test-rmse:8772.75488
[39]   test-rmse:8768.42481
[40]   test-rmse:8771.42676
```

```
[41] test-rmse:8767.30469
[42] test-rmse:8767.16504
[43] test-rmse:8757.79981
[44] test-rmse:8759.59180
[45] test-rmse:8764.60059
[46] test-rmse:8767.15234
[47] test-rmse:8771.22070
[48] test-rmse:8771.06641
```

```
[7]: print(model)
model.save_model('./xgb1.model')

model.load_model('./xgb1.model')

prediction = model.predict(xgb.DMatrix(X_test),
                           ntree_limit=model.best_ntree_limit)

plt.figure(figsize=(8, 8))
# plt.scatter(y_test[:1000], prediction[:1000], alpha=0.2)
plt.scatter(y_test, prediction, alpha=0.2)
plt.title('Evaluation between y_test=Correct Answer and Prediction. if gradient_
↳is 1, perfect')
plt.xlabel('Correct Answer',size=12)
plt.ylabel('Prediction',size=12)
plt.show()

fig, ax = plt.subplots(figsize=(8, 8))
xgb.plot_importance(model, max_num_features=12, height=0.8, ax=ax)
plt.show()
```


<xgboost.core.Booster object at 0x000001A25D6E7EB0>

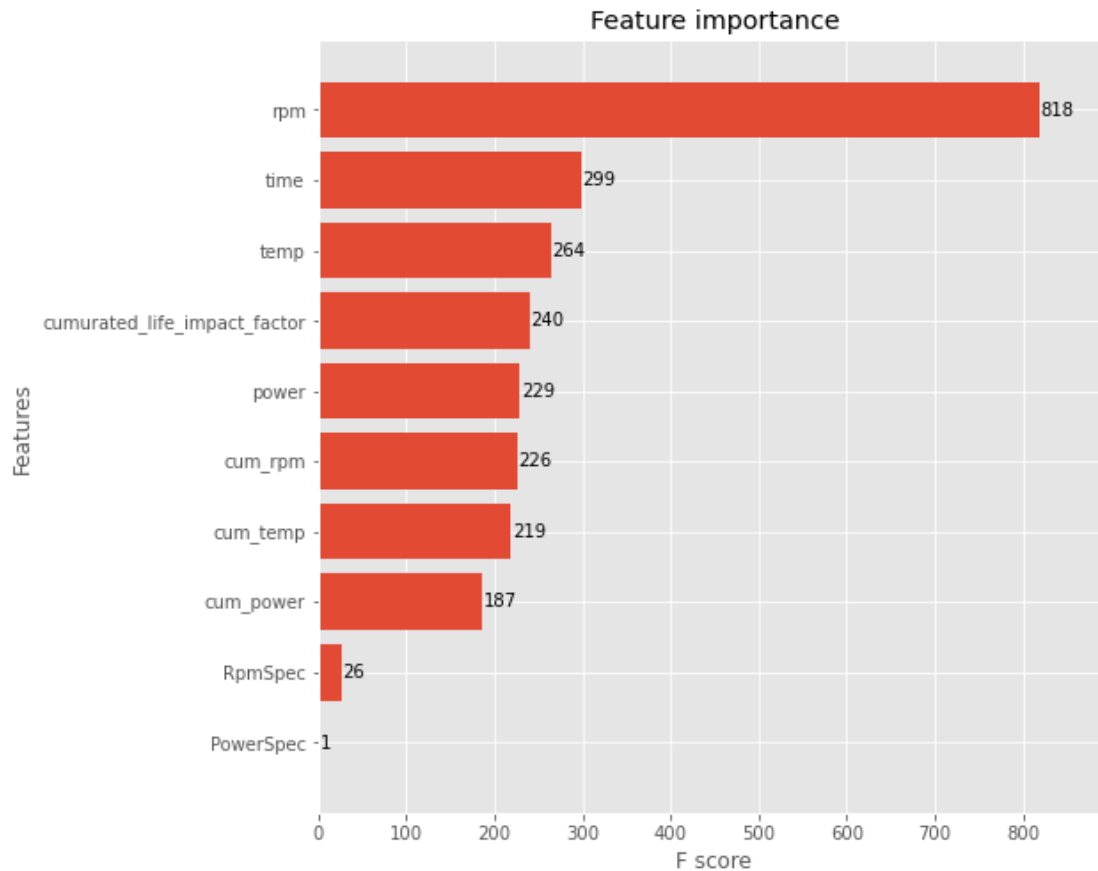
[09:07:12] WARNING: C:/Users/Administrator/workspace/xgboost-win64\_release\_1.4.0/src/objective/regression\_obj.cu:171: reg:linear is now deprecated in favor of reg:squarederror.

C:\Users\nnroc\anaconda3\lib\site-packages\xgboost\core.py:101: UserWarning: ntree\_limit is deprecated, use `iteration\_range` or model slicing instead.  
warnings.warn(



Evaluation between  $y_{\text{test}}$ =Correct Answer and Prediction. if gradient is 1, perfect





[ ]:

```
[8]: X_test_df = pd.DataFrame(X_test)
y_test_df = pd.DataFrame(y_test)
prediction_df = pd.DataFrame(prediction, columns=['remaining_life_pred'])

X_test_df_reset = X_test_df.reset_index()
y_test_df_reset = y_test_df.reset_index()
prediction_df_reset = prediction_df.reset_index()

print(y_test_df_reset)

y_test_list = y_test_df_reset['remaining_life'].to_list()
prediction_list = prediction_df_reset['remaining_life_pred'].to_list()

print(y_test_list[:5], len(y_test_list))
print(prediction_list[:5], len(prediction_list))
difference_list = []
```

```

for i in range(len(y_test_list)):
    diff = (y_test_list[i] - prediction_list[i]) / y_test_list[i]
    difference_list.append(diff)

print(difference_list[:5])

difference_df = pd.DataFrame(difference_list, columns=['defference_rate'])
print(difference_df)
print('')
print('difference_df.info()==>', difference_df.info())
print('difference_df.describe()==>', difference_df.describe())
print('length comparison==>', len(X_test_df_reset), len(y_test_df_reset),
      ↪len(prediction_df_reset), len(difference_df))

difference_df_reset = difference_df.reset_index()

plt.figure(figsize=(8, 8))
plt.hist(difference_df['defference_rate'], bins=100, range=(-1,1))
plt.title('Accuracy Verification (x=0 is correct)')
plt.show()

report_df = pd.concat([X_test_df_reset, y_test_df_reset, prediction_df_reset,
      ↪difference_df_reset], axis=1)

report_df

report_df.to_csv("./report_remaining_life.csv")

```

	index	remaining_life
0	23242	21799.93
1	12692	27539.23
2	4334	66073.56
3	11191	17004.39
4	20434	27982.34
...	...	...
4204	11936	69472.39
4205	3946	62116.93
4206	5034	11273.14
4207	1578	32261.88
4208	17770	59830.38

```

[4209 rows x 2 columns]
[21799.9262274026, 27539.23311048046, 66073.56435337233, 17004.385899645626,
27982.33607424148] 4209
[21000.203125, 31099.66015625, 59397.828125, 13450.2841796875, 29800.15625] 4209
[0.036684670125045886, -0.1292856279434508, 0.101034903954468,
0.20901088348225463, -0.06496313141746136]

```

```

      defference_rate
0          0.04
1         -0.13
2          0.10
3          0.21
4         -0.06
...
4204        0.03
4205        0.11
4206        0.31
4207        0.11
4208        0.10


```

[4209 rows x 1 columns]

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4209 entries, 0 to 4208
Data columns (total 1 columns):
#   Column          Non-Null Count  Dtype
---  -
0   defference_rate  4209 non-null   float64
dtypes: float64(1)
memory usage: 33.0 KB
difference_df.info()==> None
difference_df.describe()==>
count      4209.00
mean       -0.09
std        0.67
min       -25.61
25%       -0.04
50%        0.04
75%        0.09
max        0.62
length comparison==> 4209 4209 4209 4209

```



[ ]:

[ ]:

**Can Defects be detected?**

```
[9]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
plt.style.use('ggplot')
import seaborn as sns
```

```

df = pd.read_csv('sample_data.csv', index_col=[0])
df = df.dropna(how="any")

# print(df.head(), df.tail())
df.info()
# print(df.describe())

```

```


<class 'pandas.core.frame.DataFrame'>
Int64Index: 21041 entries, 1 to 23985
Data columns (total 19 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   sample_id                            21041 non-null  float64
1   defect                               21041 non-null  float64
2   time                                 21041 non-null  float64
3   rpm                                  21041 non-null  float64
4   temp                                 21041 non-null  float64
5   power                                21041 non-null  float64
6   cum_rpm                             21041 non-null  float64
7   cum_temp                            21041 non-null  float64
8   cum_power                           21041 non-null  float64
9   RpmSpec                             21041 non-null  float64
10  PowerSpec                           21041 non-null  float64
11  DiameterSpec                        21041 non-null  float64
12  ThicknessSpec                       21041 non-null  float64
13  QSpec                               21041 non-null  float64
14  PSpec                               21041 non-null  float64
15  LifeSpec                            21041 non-null  float64
16  cumurated_life_impact_factor        21041 non-null  float64
17  death                               21041 non-null  float64
18  remaining_life                      21041 non-null  float64
dtypes: float64(19)
memory usage: 3.2 MB

```

```

[10]: plt.figure(figsize=(8, 8))
plt.hist(df['defect'], bins=2)
plt.title('Data Histogram for prediction value')
plt.xlabel('defect (0; good, 1; defect)',size=12)
plt.show()

```



```
[11]: from sklearn.model_selection import train_test_split
import xgboost as xgb

X = df.drop(columns=['remaining_life', 'sample_id', 'defect'])
y = df['defect']

print(X.shape)
print(y.shape)

# train      test
#           20%
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=0)
```

```
print(X_train.shape)
print(X_test.shape)
```

```
(21041, 16)
(21041,)
(16832, 16)
(4209, 16)
```

```
[12]: params = {
    'silent': 1,
    'max_depth': 6,
    'min_child_weight': 1,
    'eta': 0.1,
    'tree_method': 'exact',
    'objective': 'reg:linear',
    'eval_metric': 'rmse',
    'predictor': 'cpu_predictor'
}

# GPU
# params = {
#     'silent': 1,
#     'max_depth': 6,
#     'min_child_weight': 1,
#     'eta': 0.1,
#     'tree_method': 'gpu_exact',
#     'objective': 'gpu:reg:linear',
#     'eval_metric': 'rmse',
#     'predictor': 'gpu_predictor'
# }

dtrain = xgb.DMatrix(X_train, label=y_train)
dtest = xgb.DMatrix(X_test, label=y_test)
model = xgb.train(params=params,
                  dtrain=dtrain,
                  num_boost_round=1000,
                  early_stopping_rounds=5,
                  evals=[(dtest, 'test')])
```

```
[09:07:16] WARNING: C:/Users/Administrator/workspace/xgboost-
win64_release_1.4.0/src/objective/regression_obj.cu:171: reg:linear is now
deprecated in favor of reg:squarederror.
```

```
[09:07:16] WARNING: C:/Users/Administrator/workspace/xgboost-
win64_release_1.4.0/src/learner.cc:573:
Parameters: { "silent" } might not be used.
```

This may not be accurate due to some parameters are only used in language bindings but



passed down to XGBoost core. Or some parameters are not used but slip through this verification. Please open an issue if you find above cases.

```
[0]    test-rmse:0.45762
[1]    test-rmse:0.42024
[2]    test-rmse:0.38718
[3]    test-rmse:0.35830
[4]    test-rmse:0.33291
[5]    test-rmse:0.31060
[6]    test-rmse:0.29151
[7]    test-rmse:0.27465
[8]    test-rmse:0.26037
[9]    test-rmse:0.24807
[10]   test-rmse:0.23767
[11]   test-rmse:0.22901
[12]   test-rmse:0.22181
[13]   test-rmse:0.21548
[14]   test-rmse:0.21042
[15]   test-rmse:0.20578
[16]   test-rmse:0.20226
[17]   test-rmse:0.19900
[18]   test-rmse:0.19672
[19]   test-rmse:0.19479
[20]   test-rmse:0.19304
[21]   test-rmse:0.19179
[22]   test-rmse:0.19074
[23]   test-rmse:0.18996
[24]   test-rmse:0.18924
[25]   test-rmse:0.18862
[26]   test-rmse:0.18809
[27]   test-rmse:0.18770
[28]   test-rmse:0.18738
[29]   test-rmse:0.18721
[30]   test-rmse:0.18701
[31]   test-rmse:0.18683
[32]   test-rmse:0.18686
[33]   test-rmse:0.18679
[34]   test-rmse:0.18656
[35]   test-rmse:0.18662
[36]   test-rmse:0.18658
[37]   test-rmse:0.18662
[38]   test-rmse:0.18662
```

```
[13]: print(model)
      model.save_model('./xgb1.model')
```

```

model.load_model('./xgb1.model')

prediction = model.predict(xgb.DMatrix(X_test),
                           ntree_limit=model.best_ntree_limit)

plt.figure(figsize=(8, 8))
# plt.scatter(y_test[:1000], prediction[:1000], alpha=0.2)
plt.scatter(y_test, prediction, alpha=0.2)
plt.title('Evaluation between y_test=Correct Answer and Prediction')
plt.xlabel('Correct Answer',size=12)
plt.ylabel('Prediction',size=12)
plt.show()


fig, ax = plt.subplots(figsize=(8, 8))
xgb.plot_importance(model, max_num_features=12, height=0.8, ax=ax)
plt.show()

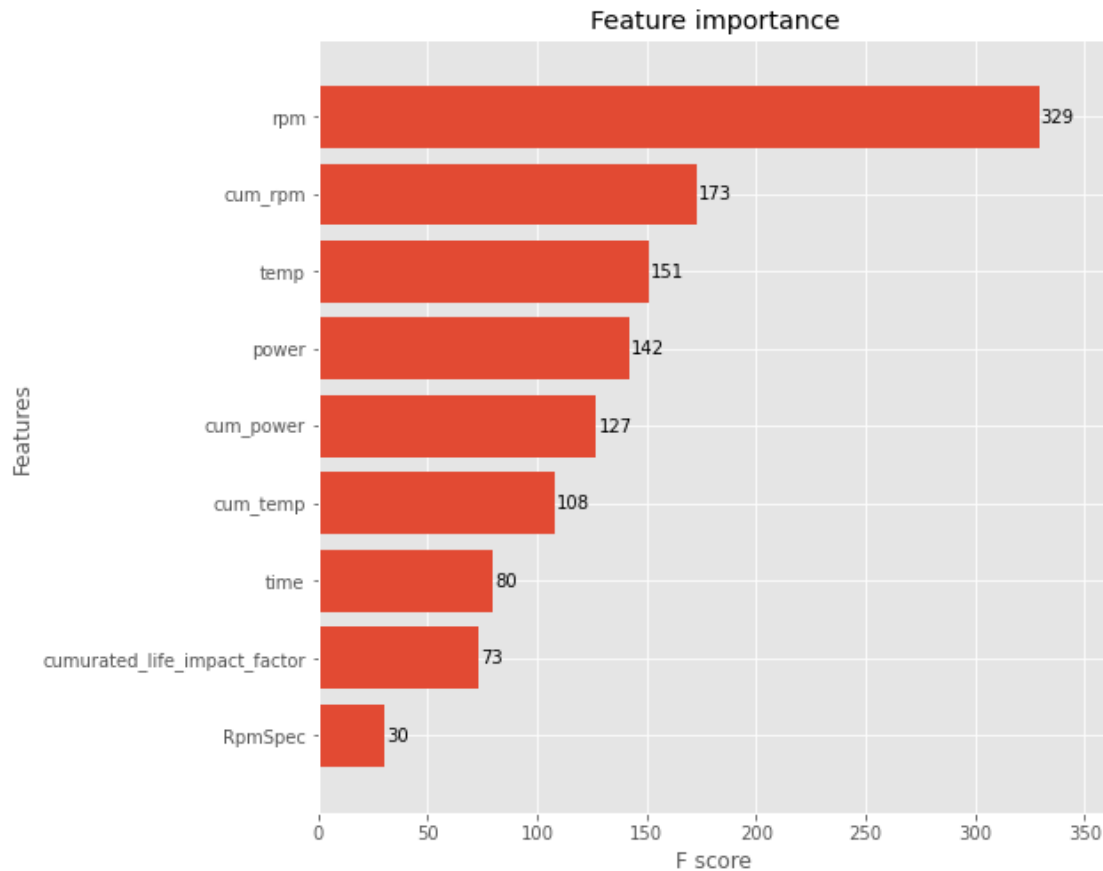
```

```

<xgboost.core.Booster object at 0x000001A25E5AE130>
[09:07:17] WARNING: C:/Users/Administrator/workspace/xgboost-
win64_release_1.4.0/src/objective/regression_obj.cu:171: reg:linear is now
deprecated in favor of reg:squarederror.
C:\Users\nnroc\anaconda3\lib\site-packages\xgboost\core.py:101: UserWarning:
ntree_limit is deprecated, use `iteration_range` or model slicing instead.
  warnings.warn(

```





```
[14]: X_test_df = pd.DataFrame(X_test)
y_test_df = pd.DataFrame(y_test)
prediction_df = pd.DataFrame(prediction, columns=['defect_pred'])

X_test_df_reset = X_test_df.reset_index()
y_test_df_reset = y_test_df.reset_index()
prediction_df_reset = prediction_df.reset_index()

print(y_test_df_reset)

y_test_list = y_test_df_reset['defect'].to_list()
prediction_list = prediction_df_reset['defect_pred'].to_list()

print(y_test_list[:5], len(y_test_list))
print(prediction_list[:5], len(prediction_list))
difference_list = []

for i in range(len(y_test_list)):
    diff = (y_test_list[i] - prediction_list[i])
```

```

        difference_list.append(diff)

print(difference_list[:5])

difference_df = pd.DataFrame(difference_list, columns=['defference_rate'])
print(difference_df)
print('')
print('difference_df.info()==>', difference_df.info())
print('difference_df.describe()==>', difference_df.describe())
print('length comparison==>', len(X_test_df_reset), len(y_test_df_reset),
      ↪len(prediction_df_reset), len(difference_df))

difference_df_reset = difference_df.reset_index()

plt.figure(figsize=(8, 8))
plt.hist(difference_df['defference_rate'], bins=100, range=(-1,1))
plt.title('Accuracy Verification (x=0 is correct)')
plt.show()

report_df = pd.concat([X_test_df_reset, y_test_df_reset, prediction_df_reset,
      ↪difference_df_reset], axis=1)

report_df

report_df.to_csv("./report_detect_defect.csv")

```

	index	defect
0	23242	0.0
1	12692	0.0
2	4334	0.0
3	11191	0.0
4	20434	0.0
...	...	...
4204	11936	0.0
4205	3946	0.0
4206	5034	0.0
4207	1578	0.0
4208	17770	0.0

[4209 rows x 2 columns]

[0.0, 0.0, 0.0, 0.0, 0.0] 4209

[0.03570520877838135, 0.013064580038189888, 0.1945408284664154, 0.013400504365563393, 0.015158231370151043] 4209

[-0.03570520877838135, -0.013064580038189888, -0.1945408284664154, -0.013400504365563393, -0.015158231370151043]

	defference_rate
0	-0.04

```

1          -0.01
2          -0.19
3          -0.01
4          -0.02
...
4204       -0.13
4205       -0.06
4206       -0.01
4207       -0.01
4208       -0.05


```

```
[4209 rows x 1 columns]
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4209 entries, 0 to 4208
Data columns (total 1 columns):
#   Column          Non-Null Count  Dtype
---  -
0   defference_rate  4209 non-null  float64
dtypes: float64(1)
memory usage: 33.0 KB
difference_df.info()==> None
difference_df.describe()==>
count      4209.00
mean        -0.01
std         0.19
min         -0.71
25%         -0.07
50%         -0.02
75%         -0.01
max         0.95
length comparison==> 4209 4209 4209 4209

```



```
[15]: print('Completed!!Completed!!!Completed!!!!Completed!!!!Completed!!!!Completed!!!
→!!Completed!!!')
```

Completed!!Completed!!!Completed!!!!Completed!!!!Completed!!!!Completed!!!!Completed!!!

```
[ ]:
```

```
[ ]:
```